

## 1. Import necessary Libraries

In [38]:

```

#For computational and random seed purpose
import numpy as np
np.random.seed(42)
#to read csv file
import pandas as pd
#To split into train and cv data
from sklearn.model_selection import train_test_split
#To compute AUROC
from sklearn.metrics import auc,roc_auc_score
#for AUROC graph
import matplotlib.pyplot as plt
#for oversampling technique
from imblearn.over_sampling import SMOTE # (https://imbalanced-learn.org/stable/references/generated/imblearn.ove
r_sampling.SMOTE.html)
#Data is imbalanced, we need calibrated model
from sklearn.calibration import CalibratedClassifierCV
#for hyperparameter tuning and Cross-validation fold
from sklearn.model_selection import GridSearchCV,StratifiedKFold,RepeatedStratifiedKFold
#to ignore the error message
import warnings
warnings.filterwarnings("ignore")
#for heatmap and other plotting technique
import seaborn as sns
#to strandize the real value data
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import LabelEncoder
#To create Knn model on datasets
from sklearn.neighbors import KNeighborsClassifier
#for roc_curve
from sklearn.metrics import roc_curve,roc_auc_score,accuracy_score

import eli5
from eli5.sklearn import PermutationImportance
import joblib
import sys
sys.modules['sklearn.externals.joblib'] = joblib
from mlxtend.feature_selection import SequentialFeatureSelector
from sklearn.linear_model import LogisticRegression
from sklearn.feature_selection import RFE
from scipy.stats import kurtosis
from scipy.stats import skew
import warnings
warnings.filterwarnings('ignore')
from catboost import CatBoostClassifier
from sklearn.preprocessing import RobustScaler

```

In [39]:

```

#locate parent directory
data_dir = "."

#Read the training data
df_train = pd.read_csv('train.csv')
print(df_train)

```

|     | id  | target | 0      | 1      | 2      | ... | 295    | 296    | 297    | 298    | 299    |
|-----|-----|--------|--------|--------|--------|-----|--------|--------|--------|--------|--------|
| 0   | 0   | 1.0    | -0.098 | 2.165  | 0.681  | ... | -2.097 | 1.051  | -0.414 | 1.038  | -1.065 |
| 1   | 1   | 0.0    | 1.081  | -0.973 | -0.383 | ... | -1.624 | -0.458 | -1.099 | -0.936 | 0.973  |
| 2   | 2   | 1.0    | -0.523 | -0.089 | -0.348 | ... | -1.165 | -1.544 | 0.004  | 0.800  | -1.211 |
| 3   | 3   | 1.0    | 0.067  | -0.021 | 0.392  | ... | 0.467  | -0.562 | -0.254 | -0.533 | 0.238  |
| 4   | 4   | 1.0    | 2.347  | -0.831 | 0.511  | ... | 1.378  | 1.246  | 1.478  | 0.428  | 0.253  |
| ... | ... | ...    | ...    | ...    | ...    | ... | ...    | ...    | ...    | ...    | ...    |
| 245 | 245 | 0.0    | -1.199 | 0.466  | -0.908 | ... | -0.243 | 0.525  | 0.281  | -0.255 | -1.136 |
| 246 | 246 | 0.0    | 0.237  | 0.233  | -0.380 | ... | 1.004  | -0.979 | 0.007  | 0.112  | -0.558 |
| 247 | 247 | 0.0    | 1.411  | -1.465 | 0.119  | ... | -0.727 | 0.461  | 0.760  | 0.168  | -0.719 |
| 248 | 248 | 1.0    | 0.620  | 1.040  | 0.184  | ... | 0.478  | -0.910 | -0.805 | 2.029  | -0.423 |
| 249 | 249 | 0.0    | 0.489  | 0.403  | 0.139  | ... | 0.812  | 0.269  | -1.454 | -0.625 | 1.474  |

[250 rows x 302 columns]

In [40]:

```
#Read test data
df_test = pd.read_csv('test.csv')
df_test
```

Out[40]:

|       | id    | 0      | 1      | 2      | 3      | 4      | 5      | 6      | 7      | 8      | 9      | 10     | 11     | 12     | 13     | 14     |      |
|-------|-------|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|------|
| 0     | 250   | 0.500  | -1.033 | -1.595 | 0.309  | -0.714 | 0.502  | 0.535  | -0.129 | -0.687 | 1.291  | 0.507  | -0.317 | 1.848  | -0.232 | -0.340 | -0.0 |
| 1     | 251   | 0.776  | 0.914  | -0.494 | 1.347  | -0.867 | 0.480  | 0.578  | -0.313 | 0.203  | 1.356  | -1.086 | 0.322  | 0.876  | -0.563 | -1.394 | 0.3  |
| 2     | 252   | 1.750  | 0.509  | -0.057 | 0.835  | -0.476 | 1.428  | -0.701 | -2.009 | -1.378 | 0.167  | -0.132 | 0.459  | -0.341 | 0.014  | 0.184  | -0.4 |
| 3     | 253   | -0.556 | -1.855 | -0.682 | 0.578  | 1.592  | 0.512  | -1.419 | 0.722  | 0.511  | 0.567  | 0.356  | -0.060 | 0.767  | -0.196 | 0.359  | 0.0  |
| 4     | 254   | 0.754  | -0.245 | 1.173  | -1.623 | 0.009  | 0.370  | 0.781  | -1.763 | -1.432 | -0.930 | -0.098 | 0.896  | 0.293  | -0.259 | 0.030  | -0.6 |
| ...   | ...   | ...    | ...    | ...    | ...    | ...    | ...    | ...    | ...    | ...    | ...    | ...    | ...    | ...    | ...    | ...    | ...  |
| 19745 | 19995 | 1.069  | 0.517  | -0.690 | 0.241  | 0.913  | -0.859 | 0.093  | -0.359 | -0.047 | 0.713  | 2.191  | 0.774  | -0.110 | -0.721 | 0.375  | 0.5  |
| 19746 | 19996 | -0.529 | 0.438  | 0.672  | 1.436  | -0.720 | 0.698  | -0.350 | 2.150  | -1.241 | -0.167 | -0.188 | 0.541  | -0.392 | 1.727  | -0.965 | 0.5  |
| 19747 | 19997 | -0.554 | -0.936 | -1.427 | 0.027  | -0.539 | 0.994  | -1.832 | -1.156 | 0.474  | 1.483  | 1.524  | 0.143  | -0.607 | -1.142 | 2.786  | -0.3 |
| 19748 | 19998 | -0.746 | 1.205  | 0.750  | -0.236 | 1.139  | -1.727 | -0.677 | -1.254 | -0.099 | -0.724 | 0.014  | -0.575 | -0.142 | 1.171  | -0.198 | 0.3  |
| 19749 | 19999 | 0.736  | -0.216 | -0.110 | -1.404 | -0.265 | -1.770 | 0.715  | 0.469  | 1.077  | 0.333  | -0.994 | -0.331 | 1.009  | 0.607  | -1.729 | 1.4  |

19750 rows × 301 columns



In [41]:

```
df_train.dropna(inplace=True)
df_test.dropna(inplace=True)
```

## Observation

We have read the training and test dataset. After reading both of dataset, we got it know that test dataset is having more features compare to training dataset

In [42]:

```
def feature_engg(df,test=False):
    """
    perform feature Engineering in basic statistics, trignometry, hyperbolic and exponential function

    parameters:
    """
    if test:
        data = df.drop(['id'],axis=1)
    else:
        data = df.drop(['id','target'],axis=1)

    #mean and std
    df['mean'] = np.mean(data,axis=1) # taking mean value along with column
    df['std'] = np.std(data,axis=1) # taking std along with column
    df['median'] = np.median(data,axis=1)
    df['min'] = np.min(data,axis=1)
    df['max'] = np.max(data,axis=1)

    # applying trignometric function
    df['sin_mean'] = np.sin(df['mean'])
    df['cos_mean'] = np.cos(df['mean'])
    df['tan_mean'] = np.tan(df['mean'])
    df['sin_std'] = np.sin(df['std'])
    df['cos_std'] = np.cos(df['std'])
    df['tan_std'] = np.tan(df['std'])

    df['sin_median'] = np.sin(df['median'])
    df['cos_median'] = np.cos(df['median'])
    df['tan_median'] = np.tan(df['median'])
    sin_data = np.sin(data) #calculated the sin_data
    cos_data = np.cos(data) #calculated the cos_data
    tan_data = np.tan(data) #calculated the tan_data

    df['mean_sin'] = np.mean(sin_data,axis=1) #calculating the mean of sin_data
    df['cos_mean'] = np.mean(cos_data,axis=1) #calculating the mean of cos_data
    df['tan_mean'] = np.mean(tan_data,axis=1) #calculating the mean of tan_data
```

```

df['mean_cos'] = np.mean(cos_data,axis=1) #calculating the mean of cos_data

df['mean_tan'] = np.mean(tan_data,axis=1) #calculating the mean of tan_data

#hyperbolic function

sinh_data = np.sinh(data)
cosh_data = np.cosh(data)
tanh_data = np.tanh(data)
arcsinh_data = np.arcsinh(data)
arccosh_data = np.arccosh(data)

df['mean_sinh'] = np.mean(sinh_data,axis=1)
df['mean_cosh'] = np.mean(cosh_data,axis=1)
df['mean_tanh'] = np.mean(tanh_data,axis=1)
df['mean_arsinh'] = np.mean(arcsinh_data,axis=1)
df['mean_arcosh'] = np.mean(arccosh_data,axis=1)
df['sinh_mean'] = np.sinh(df['mean'])

df['tanh_mean'] = np.tanh(df['mean'])
df['arsinh_mean'] = np.arcsinh(df['mean'])
df['sinh_std'] = np.sinh(df['std'])
df['cosh_std'] = np.cosh(df['std'])
df['tanh_std'] = np.tanh(df['std'])
df['sinh_median'] = np.sinh(df['median'])
df['cosh_median'] = np.cosh(df['median'])
df['tanh_median'] = np.tanh(df['median'])

#exponential function

exp_data = np.exp(data)
expm1_data = np.expm1(data)
exp2_data = np.exp2(data)

df['mean_exp'] = np.mean(exp_data,axis=1)
df['mean_expm1'] = np.mean(expm1_data,axis=1)
df['mean_exp2'] = np.mean(exp2_data,axis=1)
df['exp1_mean'] = np.exp(df['mean'])
df['expm1_mean'] = np.expm1(df['mean'])
df['exp2_mean'] = np.exp2(df['mean'])
df['exp1_median'] = np.exp(df['median'])
df['expm1_median'] = np.expm1(df['median'])
df['exp2_median'] = np.exp2(df['median'])

df['exp1_std'] = np.exp(df['std'])
df['expm1_std'] = np.expm1(df['std'])
df['exp2_std'] = np.exp2(df['std'])
# Polynomial FE
# X**2
df['mean_x2'] = np.mean(np.power(data,2), axis=1)
# X**3
df['mean_x3'] = np.mean(np.power(data,3), axis=1)
# X**4
df['mean_x4'] = np.mean(np.power(data,4), axis=1)
# X**5
df['mean_x5'] = np.mean(np.power(data,5), axis=1)
# X**6
df['mean_x6'] = np.mean(np.power(data,6), axis=1)
# X**7
df['mean_x7'] = np.mean(np.power(data,7), axis=1)
#logithm FE
df['x2_mean'] = np.power(df['mean'],2)
# X**3
df['x3_mean'] = np.power(df['mean'],3)
# X**4
df['x4_mean'] = np.power(df['mean'],4)
# X**5
df['x5_mean'] = np.power(df['mean'],5)
# X**6
df['x6_mean'] = np.power(df['mean'],6)
# X**7
df['x7_mean'] = np.power(df['mean'],7)
#skewness and kurtosis
skew_data = skew(data)
kurtosis_data = kurtosis(data)
df['skewness'] = np.mean(skew_data)

df['kurtosis'] = np.mean(kurtosis_data)
data['mean_skewness'] = skew(df['mean'])
data['mean_kurtosis'] = kurtosis(df['mean'])
df['x2_median'] = np.power(df['median'],2)
# X**3
df['x3_median'] = np.power(df['median'],3)
# X**4
df['x4_median'] = np.power(df['median'],4)
# X**5
df['x5_median'] = np.power(df['median'],5)
# X**6
df['x6_median'] = np.power(df['median'],6)
# X**7
df['x7_median'] = np.power(df['median'],7)

```

```
# X**4
df['x4_median'] = np.power(df['median'],4)
# X**5
df['x5_median'] = np.power(df['median'],5)
# X**6
df['x6_median'] = np.power(df['median'],6)
# X**7
df['x7_median'] = np.power(df['median'],7)

return df
```

In [43]:

```
df_train = feature_engg(df_train)
df_train.head(5)
```

Out[43]:

|   | id | target | 0      | 1      | 2      | 3      | 4      | 5      | 6      | 7      | 8      | 9      | 10     | 11     | 12     | 13     | 14     |     |
|---|----|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|-----|
| 0 | 0  | 1.0    | -0.098 | 2.165  | 0.681  | -0.614 | 1.309  | -0.455 | -0.236 | 0.276  | -2.246 | 1.825  | -0.912 | -0.107 | 0.305  | 0.102  | 0.826  | 0.4 |
| 1 | 1  | 0.0    | 1.081  | -0.973 | -0.383 | 0.326  | -0.428 | 0.317  | 1.172  | 0.352  | 0.004  | -0.291 | 2.907  | 1.085  | 2.144  | 1.540  | 0.584  | 1.1 |
| 2 | 2  | 1.0    | -0.523 | -0.089 | -0.348 | 0.148  | -0.022 | 0.404  | -0.023 | -0.172 | 0.137  | 0.183  | 0.459  | 0.478  | -0.425 | 0.352  | 1.095  | 0.3 |
| 3 | 3  | 1.0    | 0.067  | -0.021 | 0.392  | -1.637 | -0.446 | -0.725 | -1.035 | 0.834  | 0.503  | 0.274  | 0.335  | -1.148 | 0.067  | -1.010 | 1.048  | -1. |
| 4 | 4  | 1.0    | 2.347  | -0.831 | 0.511  | -0.021 | 1.225  | 1.594  | 0.585  | 1.509  | -0.012 | 2.198  | 0.190  | 0.453  | 0.494  | 1.478  | -1.412 | 0.2 |

5 rows x 365 columns

In [44]:

```
df_test = feature_engg(df_test,True)
df_test.head(5)
```

Out[44]:

|   | id  | 0      | 1      | 2      | 3      | 4      | 5     | 6      | 7      | 8      | 9      | 10     | 11     | 12     | 13     | 14     | 15     |     |
|---|-----|--------|--------|--------|--------|--------|-------|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|-----|
| 0 | 250 | 0.500  | -1.033 | -1.595 | 0.309  | -0.714 | 0.502 | 0.535  | -0.129 | -0.687 | 1.291  | 0.507  | -0.317 | 1.848  | -0.232 | -0.340 | -0.051 | 0.  |
| 1 | 251 | 0.776  | 0.914  | -0.494 | 1.347  | -0.867 | 0.480 | 0.578  | -0.313 | 0.203  | 1.356  | -1.086 | 0.322  | 0.876  | -0.563 | -1.394 | 0.385  | 1.  |
| 2 | 252 | 1.750  | 0.509  | -0.057 | 0.835  | -0.476 | 1.428 | -0.701 | -2.009 | -1.378 | 0.167  | -0.132 | 0.459  | -0.341 | 0.014  | 0.184  | -0.460 | -0. |
| 3 | 253 | -0.556 | -1.855 | -0.682 | 0.578  | 1.592  | 0.512 | -1.419 | 0.722  | 0.511  | 0.567  | 0.356  | -0.060 | 0.767  | -0.196 | 0.359  | 0.080  | -0. |
| 4 | 254 | 0.754  | -0.245 | 1.173  | -1.623 | 0.009  | 0.370 | 0.781  | -1.763 | -1.432 | -0.930 | -0.098 | 0.896  | 0.293  | -0.259 | 0.030  | -0.661 | 0.  |

5 rows x 364 columns

Observation

We applied feature engineering on the training dataset and came up with new features

In [45]:

```
X_train = (df_train.drop(['id','target'],axis = 1))
X_test = (df_test.drop(['id'],axis = 1))

y_train = df_train['target']

smote = SMOTE()
X_train,y_train = smote.fit_resample(X_train,y_train)
```

In [46]:

```
X_train.shape
```

Out[46]:

```
(320, 363)
```

## Observation

We have dropped the labeled data from both training and test dataset and applied Oversampling Technique

In [47]:

```
stand = StandardScaler()  
X_train = stand.fit_transform(X_train)  
X_test = stand.transform(X_test)
```

In [48]:

```
X_train.shape
```

Out[48]:

```
(320, 363)
```

In [49]:

```
X_train.shape
```

Out[49]:

```
(320, 363)
```

## Observation

We have used StandardScaler() method to standardize the both training and test dataset.

### Used GridSearch for hyper-parameter tuning

In [50]:

```
def hyperparameter_model(models, params):  
    '''  
    Hyperparameter tuning with StratifiedKFold follow by GridSearchCV follow by_  
    ,→CalibratedClassifier  
    Parameters:  
    models: Instance of the model  
    params: list of parameters with value fr tuning (dict)  
    Return:  
    grid_clf: return gridsearch model  
    '''  
    # Perform KCrossValidation with stratified target  
    str_cv = StratifiedKFold(n_splits=11, random_state=42,shuffle=True)  
    # Perform Hyperparamter using GridSearchCV  
    grid_clf = GridSearchCV(models, params, cv=str_cv, return_train_score=True,scoring='roc_auc')  
    # Fit the train model to evaluate score  
    grid_clf.fit(X_train, y_train)  
    return grid_clf
```

In [14]:

```
#kNN (See Docs: https://scikit-learn.org/stable/modules/generated/sklearn.neighbors.KNeighborsClassifier.html)  
# List of params  
# List of params  
params = {'n_neighbors':np.arange(3,51,2).tolist(), 'algorithm': ['kd_tree','brute']}  
  
# Instance of knn model  
knn_model = KNeighborsClassifier()  
# Call hyperparameter for find the best params as possible  
knn_clf = hyperparameter_model(knn_model, params)
```

In [16]:

```
print(knn_clf.best_params_)
```

```
{'algorithm': 'kd_tree', 'n_neighbors': 39}
```

```
knn_model = KNeighborsClassifier(**knn_clf.best_params_)
knn_model.fit(X_train,y_train)
```

```
KNeighborsClassifier(algorithm='kd_tree', n_neighbors=39)
```

```
y_pred = knn_model.predict(X_train)
print(y_pred)
```

```
In [19]:  
  
train_auc = roc_auc_score(y_train,y_pred)  
print(train_auc)
```

```
In [20]:  
y_predict = knn_model.predict_proba(X_test)[: ,1]
```

In [21]:

```
y_pred_lr_test = pd.DataFrame({"ID": df_test['id'], "Target": y_predict})

y_pred_lr_test.to_csv('submission_knn1.csv', index=False)
y_pred_lr_test.head(20)
```

Out[21]:

|    | ID  | Target   |
|----|-----|----------|
| 0  | 250 | 0.333333 |
| 1  | 251 | 0.051282 |
| 2  | 252 | 0.051282 |
| 3  | 253 | 0.025641 |
| 4  | 254 | 0.128205 |
| 5  | 255 | 0.051282 |
| 6  | 256 | 0.025641 |
| 7  | 257 | 0.025641 |
| 8  | 258 | 0.102564 |
| 9  | 259 | 0.102564 |
| 10 | 260 | 0.076923 |
| 11 | 261 | 0.128205 |
| 12 | 262 | 0.025641 |
| 13 | 263 | 0.102564 |
| 14 | 264 | 0.051282 |
| 15 | 265 | 0.076923 |
| 16 | 266 | 0.128205 |
| 17 | 267 | 0.076923 |
| 18 | 268 | 0.025641 |
| 19 | 269 | 0.076923 |

NameSubmittedWait timeExecution timeScore

submission\_knn1.csvjust now1 seconds0 seconds0.555

Complete

[Jump to your position on the leaderboard](#)

test\_auc = 0.555

Observation

As the first model, We have used KNN algorithm to trained the model with best parameter (algorithm='kd\_tree',n\_neighbours = 39) and got training AUC = 0.50 and Test\_AUC = 0.55. It is less accurate but model is not overfitted.

Logistic Regresstion Applied

In [51]:

```
#ref= https://scikit-learn.org/stable/modules/generated/sklearn.linear\_model.LogisticRegression.html

params = {'penalty':['l1','l2','elasticnet'],'C':[10**i for i in range(-4,5)], 'solver':['liblinear','sag','saga'
]}

#the instance of Logistic Regression

log_model = LogisticRegression(random_state=42)

#Call Hyper-parameter function to get best hyperparameter tuning

log_clf = hyperparameter_model(log_model,params)
```

In [52]:

```
print(log_clf.best_params_)

{'C': 1, 'penalty': 'l1', 'solver': 'saga'}
```

In [53]:

```
from sklearn import linear_model

model = LogisticRegression(penalty='l1', C=1, solver='saga')

model.fit(X_train,y_train)
```

Out[53]:

```
LogisticRegression(C=1, penalty='l1', solver='saga')
```

In [54]:

```
y_pred = model.predict(X_train)
print(y_pred)
```

```
[1. 0. 1. 1. 1. 1. 1. 1. 1. 1. 0. 1. 1. 0. 0. 1. 1. 0. 1. 1. 1. 0. 1.
 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 0. 1. 1. 1. 1. 1. 1.
 0. 0. 1. 1. 0. 1. 0. 1. 0. 1. 0. 0. 1. 1. 0. 0. 0. 1. 0. 1. 1. 1. 0. 1.
 1. 1. 1. 0. 1. 1. 0. 1. 1. 0. 0. 1. 1. 1. 0. 1. 1. 1. 1. 0. 1. 1. 1. 1.
 0. 1. 1. 1. 1. 0. 0. 1. 0. 1. 1. 0. 0. 1. 1. 1. 1. 1. 0. 1. 1. 0. 0. 1.
 0. 1. 0. 1. 0. 1. 0. 0. 0. 1. 1. 1. 1. 1. 0. 1. 0. 1. 1. 0. 1. 0. 1. 0.
 1. 1. 1. 0. 1. 0. 1. 1. 0. 0. 1. 1. 0. 0. 1. 1. 0. 1. 1. 0. 1. 1. 1.
 0. 0. 0. 0. 0. 1. 1. 1. 0. 0. 1. 1. 0. 0. 0. 1. 1. 1. 1. 1. 0. 1. 1.
 1. 0. 0. 1. 1. 1. 0. 1. 1. 0. 1. 1. 1. 0. 1. 0. 1. 0. 0. 0. 0. 1. 1. 0.
 1. 1. 0. 1. 1. 1. 1. 1. 0. 1. 0. 0. 1. 1. 1. 1. 0. 0. 1. 1. 1. 1. 0.
 0. 1. 1. 0. 0. 0. 0. 0. 1. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0. 0. 0.]
```

In [55]:

```
train_auc_lr = roc_auc_score(y_train,y_pred)
print(train_auc_lr)
```

```
1.0
```

In [56]:

```
y_pred_lr_test = model.predict_proba(X_test)[: ,1]
print(y_pred_lr_test)
```

```
[0.0296196  0.33341024 0.86344733 ... 0.87014075 0.99319125 0.17739898]
```



In [57]:

```
y_pred_lr_test = pd.DataFrame({"ID": df_test['id'], "Target": y_pred_lr_test})
```

```
y_pred_lr_test.to_csv('submission_logs1.csv', index=False)  
y_pred_lr_test.head(20)
```

Out[57]:

|    | ID  | Target   |
|----|-----|----------|
| 0  | 250 | 0.029620 |
| 1  | 251 | 0.333410 |
| 2  | 252 | 0.863447 |
| 3  | 253 | 0.997246 |
| 4  | 254 | 0.644034 |
| 5  | 255 | 0.386384 |
| 6  | 256 | 0.197760 |
| 7  | 257 | 0.420377 |
| 8  | 258 | 0.975555 |
| 9  | 259 | 0.187157 |
| 10 | 260 | 0.601244 |
| 11 | 261 | 0.171481 |
| 12 | 262 | 0.039011 |
| 13 | 263 | 0.879241 |
| 14 | 264 | 0.529214 |
| 15 | 265 | 0.940078 |
| 16 | 266 | 0.966850 |
| 17 | 267 | 0.685696 |
| 18 | 268 | 0.309618 |
| 19 | 269 | 0.760107 |

Your most recent submission

| Name                 | Submitted | Wait time | Execution time | Score |
|----------------------|-----------|-----------|----------------|-------|
| submission_logs1.csv | just now  | 1 seconds | 0 seconds      | 0.773 |

Complete

[Jump to your position on the leaderboard](#) ▼

logistic\_test\_auc = 0.773

## Observation

We have used Logistic Regression algorithm to trained the model with best parameter ( $c=1$ ,  $\text{penalty}=l1$ ,  $\text{solver}=saga$ ) and got training AUC = 1.0 and Test\_AUC = 0.773. Training AUC is good and test auc litte bit down compare to training AUC. But it is decent.

## Support Vector Machine

In [58]:

```
from sklearn.svm import SVC
```

In [60]:

```
#ref = https://scikit-learn.org/stable/modules/svm.html

params = {'C':[10**i for i in range(-4,5)], 'kernel':['linear','poly','sigmoid','rbf']}

#The instance of SVC

svc_model = SVC(random_state=42)
#call the hyper-parameter function to get best parameters

svc_clf = hyperparameter_model(svc_model,params)
```

In [61]:

```
print(svc_clf.best_params_)

{'C': 1, 'kernel': 'poly'}
```

In [62]:

```
svc_clf = SVC(C = 1, kernel = 'poly',probability=True)
svc_clf.fit(X_train,y_train)
```

Out[62]:

```
SVC(C=1, kernel='poly', probability=True)
```

In [63]:

```
y_pred = svc_clf.predict(X_train)
train_svm_auc = roc_auc_score(y_train,y_pred)
print(train_svm_auc)
```

```
1.0
```

In [64]:

```
y_pred_svc_test = svc_clf.predict_proba(X_test)[:,-1]
```

In [65]:

```
y_pred_svc_test = pd.DataFrame({"ID": df_test['id'], "Target": y_pred_svc_test})
```

```
y_pred_svc_test.to_csv('submission_svm1.csv', index=False)  
y_pred_svc_test.head(20)
```

Out[65]:

|    | ID  | Target   |
|----|-----|----------|
| 0  | 250 | 0.010239 |
| 1  | 251 | 0.799375 |
| 2  | 252 | 0.782953 |
| 3  | 253 | 0.776982 |
| 4  | 254 | 0.719545 |
| 5  | 255 | 0.762792 |
| 6  | 256 | 0.761592 |
| 7  | 257 | 0.770146 |
| 8  | 258 | 0.851095 |
| 9  | 259 | 0.745677 |
| 10 | 260 | 0.765826 |
| 11 | 261 | 0.558469 |
| 12 | 262 | 0.722451 |
| 13 | 263 | 0.758217 |
| 14 | 264 | 0.755097 |
| 15 | 265 | 0.784517 |
| 16 | 266 | 0.867675 |
| 17 | 267 | 0.792677 |
| 18 | 268 | 0.729581 |
| 19 | 269 | 0.715112 |

| Name                | Submitted | Wait time | Execution time | Score |
|---------------------|-----------|-----------|----------------|-------|
| submission_svm1.csv | just now  | 1 seconds | 0 seconds      | 0.644 |

Complete

[Jump to your position on the leaderboard](#) ▼

Test\_SVM\_auc = 0.644

## Observation

We have used Support Vector algorithm to trained the model with best parameter ('C': 1, 'kernel': 'Poly') and got training AUC = 1.0 and Test\_AUC = 0.64. Training AUC is good and test auc a bit down compare to training AUC. But it is decent.

## Ensemble Model : Random Forest

In [66]:

```
from sklearn.ensemble import RandomForestClassifier
```

In [67]:

```
#https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestClassifier.html
params = {'n_estimators': [10,20,30,40,50,60,100,200,300,400,500], 'max_depth': [2,3,5,7,9]}

#The instance of model

rdf_model = RandomForestClassifier(random_state=42)

# Call the hyperparameter function to get best parameter
rdf_clf = hyperparameter_model(rdf_model,params)
```

In [68]:

```
print(rdf_clf.best_params_)

{'max_depth': 5, 'n_estimators': 400}
```

In [69]:

```
rdf_clf = RandomForestClassifier(**rdf_clf.best_params_)
rdf_clf.fit(X_train,y_train)
```

Out[69]:

```
RandomForestClassifier(max_depth=5, n_estimators=400)
```

In [70]:

```
y_pred = rdf_clf.predict(X_train)
train_rdf_auc = roc_auc_score(y_train,y_pred)
print(train_rdf_auc)
```

```
1.0
```

In [71]:

```
y_pred_rdf_test = rdf_clf.predict_proba(X_test)[: ,1]
```

In [72]:

```
y_pred_rdf_test = pd.DataFrame({"ID": df_test['id'], "Target": y_pred_rdf_test})

y_pred_rdf_test.to_csv('submission_rdf1.csv', index=False)
y_pred_rdf_test.head(20)
```

Out[72]:

|    | ID  | Target   |
|----|-----|----------|
| 0  | 250 | 0.516006 |
| 1  | 251 | 0.568699 |
| 2  | 252 | 0.575418 |
| 3  | 253 | 0.579637 |
| 4  | 254 | 0.535900 |
| 5  | 255 | 0.507414 |
| 6  | 256 | 0.512040 |
| 7  | 257 | 0.527653 |
| 8  | 258 | 0.630538 |
| 9  | 259 | 0.463896 |
| 10 | 260 | 0.604303 |
| 11 | 261 | 0.430581 |
| 12 | 262 | 0.465354 |
| 13 | 263 | 0.519210 |
| 14 | 264 | 0.471482 |
| 15 | 265 | 0.606225 |
| 16 | 266 | 0.621462 |
| 17 | 267 | 0.566494 |
| 18 | 268 | 0.514597 |
| 19 | 269 | 0.463538 |

Your most recent submission

| Name                | Submitted | Wait time | Execution time | Score |
|---------------------|-----------|-----------|----------------|-------|
| submission_rdf1.csv | just now  | 1 seconds | 1 seconds      | 0.715 |

Complete

[Jump to your position on the leaderboard](#) ▼

Test\_rdf\_auc : 0.715

Observation

We have used Random Forest algorithm to trained the model with best parameter ('max\_depth': 5, 'n\_estimators': 400) and got training AUC = 1.0 and Test\_AUC = 0.72. Model is accurate and it is not overfitted.

Decision Tree Classifier

In [95]:

```
from sklearn.tree import DecisionTreeClassifier
```

In [96]:

```
#ref =https://scikit-learn.org/stable/modules/generated/sklearn.tree.DecisionTreeClassifier.html
params = {'max_depth':[2,3,5,7,9], 'criterion':['gini', 'entropy'], 'min_samples_split':[2,3,4,5,6,]}
#The instance of Decision Tree Classifier
tree_model = DecisionTreeClassifier(random_state=42)
#Call Hyperparameter function to get best parameter
tree_clf = hyperparameter_model(tree_model, params)
```

In [97]:

```
print(tree_clf.best_params_)
{'criterion': 'gini', 'max_depth': 9, 'min_samples_split': 6}
```

In [98]:

```
tree_clf = DecisionTreeClassifier(**tree_clf.best_params_)
tree_clf.fit(X_train, y_train)
```

Out[98]:

```
DecisionTreeClassifier(max_depth=9, min_samples_split=6)
```

In [99]:

```
y_pred = tree_clf.predict(X_train)
train_tree_auc = roc_auc_score(y_train, y_pred)
print(train_tree_auc)
```

0.99375

In [100]:

```
y_pred_tree_test = tree_clf.predict_proba(X_test)[: , 1]
```

In [101]:

```
y_pred_tree_test = pd.DataFrame({"ID": df_test['id'], "Target": y_pred_tree_test})

y_pred_tree_test.to_csv('submission_tree1.csv', index=False)
y_pred_tree_test.head(20)
```

Out[101]:

|    | ID  | Target   |
|----|-----|----------|
| 0  | 250 | 1.000000 |
| 1  | 251 | 0.000000 |
| 2  | 252 | 0.000000 |
| 3  | 253 | 1.000000 |
| 4  | 254 | 0.000000 |
| 5  | 255 | 0.000000 |
| 6  | 256 | 1.000000 |
| 7  | 257 | 1.000000 |
| 8  | 258 | 1.000000 |
| 9  | 259 | 0.000000 |
| 10 | 260 | 0.000000 |
| 11 | 261 | 0.000000 |
| 12 | 262 | 0.333333 |
| 13 | 263 | 1.000000 |
| 14 | 264 | 0.000000 |
| 15 | 265 | 1.000000 |
| 16 | 266 | 1.000000 |
| 17 | 267 | 1.000000 |
| 18 | 268 | 0.000000 |
| 19 | 269 | 1.000000 |

| Name                 | Submitted | Wait time | Execution time | Score |
|----------------------|-----------|-----------|----------------|-------|
| submission_tree1.csv | just now  | 1 seconds | 0 seconds      | 0.599 |

Complete

[Jump to your position on the leaderboard](#) ▼

Observation

Decision Tree is working poorly in this case

XGBoost Classifier

In [102]:

```
from xgboost import XGBClassifier
```

In [103]:

```
#list of hyper-parameter

params = {'max_depth':[2,3,5,7,9], 'n_estimators':[10,20,30,40,50,100,200,400,500]}

# The instance of XGBClassifier

xg_model = XGBClassifier()
# call hyperparameter function to get best parameter

xg_clf = hyperparameter_model(xg_model,params)
```

In [104]:

```
print(xg_clf.best_params_)

{'max_depth': 2, 'n_estimators': 400}
```

In [105]:

```
xg_clf = XGBClassifier(**xg_clf.best_params_)
xg_clf.fit(X_train,y_train)
```

Out[105]:

```
XGBClassifier(max_depth=2, n_estimators=400)
```

In [107]:

```
y_pred = xg_clf.predict(X_train)
```

In [108]:

```
train_xgboost_auc = roc_auc_score(y_train,y_pred)
print(train_xgboost_auc)
```

1.0

In [109]:

```
y_pred_xg_test = xg_clf.predict_proba(X_test)[:,:1]
print(y_pred_xg_test)
```

```
[0.7702675  0.4989509  0.2832494  ... 0.35014024 0.9550946  0.3850446 ]
```

In [110]:

```
y_pred_xg_test = pd.DataFrame({"ID": df_test['id'], "Target": y_pred_xg_test})
```

```
y_pred_xg_test.to_csv('submission_xgboost1.csv', index=False)
y_pred_xg_test.head(10)
```

Out[110]:

|   | ID  | Target   |
|---|-----|----------|
| 0 | 250 | 0.770267 |
| 1 | 251 | 0.498951 |
| 2 | 252 | 0.283249 |
| 3 | 253 | 0.995090 |
| 4 | 254 | 0.698269 |
| 5 | 255 | 0.321344 |
| 6 | 256 | 0.348547 |
| 7 | 257 | 0.270397 |
| 8 | 258 | 0.997484 |
| 9 | 259 | 0.552730 |



| Name                    | Submitted | Wait time | Execution time | Score |
|-------------------------|-----------|-----------|----------------|-------|
| submission_xgboost1.csv | just now  | 1 seconds | 0 seconds      | 0.784 |

Complete

[Jump to your position on the leaderboard ▾](#)

xgboost\_test\_auc = 0.784

## Observation

We have used Xgboost algorithm to trained the model with best parameter ('max\_depth': 2, 'n\_estimators': 400) and got training AUC = 1.0 and Test\_AUC = 0.78. Model is accurate and it is not overfitted.

## Stacking Classifier

In [111]:

```
import six
import sys
sys.modules['sklearn.externals.six'] = six
```

In [112]:

```
from mlxtend.classifier import StackingClassifier
```

In [113]:

```
#classifier 1
knn_model = KNeighborsClassifier(algorithm='kd_tree',n_neighbors = 39)
knn_model.fit(X_train,y_train)

#Classifier 2
model = LogisticRegression(penalty='l1', C=1, solver='saga')
model.fit(X_train,y_train)

#Classifier 3
svc_clf = SVC(C =1, kernel = 'poly',probability=True)
svc_clf.fit(X_train,y_train)

#classifier 3
rdf_clf = RandomForestClassifier(max_depth=5, n_estimators=400)
rdf_clf.fit(X_train,y_train)
#classifier 4

tree_clf = DecisionTreeClassifier(criterion = 'gini', max_depth = 9, min_samples_split = 6)
tree_clf.fit(X_train,y_train)

#classifier 5
xg_clf = XGBClassifier(max_depth = 2, n_estimators = 400)
xg_clf.fit(X_train,y_train)

#Stacking Classifier

sclf = StackingClassifier(classifiers=[knn_model,model,svc_clf,rdf_clf,tree_clf,xg_clf],meta_classifier=model,use
_probab=True)

#fit the model
sclf.fit(X_train,y_train)

#predict in probabilities

y_pred = sclf.predict(X_train)
```

In [114]:

```
train_auc = roc_auc_score(y_train,y_pred)
print(train_auc)
```

1.0

In [115]:

```
y_pred_stack_test = sclf.predict_proba(X_test)[: ,1]
```

In [116]:

```
y_pred_stack_test = pd.DataFrame({"ID": df_test['id'], "Target": y_pred_stack_test})

y_pred_stack_test.to_csv('submission_stack1.csv', index=False)
y_pred_stack_test.head(20)
```

Out[116]:

|    | ID  | Target   |
|----|-----|----------|
| 0  | 250 | 0.384830 |
| 1  | 251 | 0.664552 |
| 2  | 252 | 0.444153 |
| 3  | 253 | 0.984780 |
| 4  | 254 | 0.739978 |
| 5  | 255 | 0.819734 |
| 6  | 256 | 0.833805 |
| 7  | 257 | 0.795113 |
| 8  | 258 | 0.989374 |
| 9  | 259 | 0.650674 |
| 10 | 260 | 0.886006 |
| 11 | 261 | 0.119933 |
| 12 | 262 | 0.340979 |
| 13 | 263 | 0.982276 |
| 14 | 264 | 0.264743 |
| 15 | 265 | 0.984879 |
| 16 | 266 | 0.712580 |
| 17 | 267 | 0.982666 |
| 18 | 268 | 0.381315 |
| 19 | 269 | 0.941461 |

Your most recent submission

| Name                  | Submitted | Wait time | Execution time | Score |
|-----------------------|-----------|-----------|----------------|-------|
| submission_stack1.csv | just now  | 1 seconds | 0 seconds      | 0.759 |

Complete

[Jump to your position on the leaderboard](#) ▼

Summary of All Models

In [1]:

```
from texttable import Texttable
t = Texttable()
t.add_rows([[ 'Model', 'Hyper-parameter', 'Train AUC', 'Test AUC'], [ 'Knn_Model', r"(algorithm='kd_tree', n_neighbors=39)", 0.50, 0.55],
            [ 'logistic Regresstion', r"{'C': 1, 'penalty': 'l1', 'solver': 'saga'}", 1.0, 0.77], [ 'Support Vector Machine',
            r"{'C': 1, 'kernel': 'poly'}", 1.0, 0.64], [ 'XGBoost Classifier', r"{'max_depth': 2, 'n_estimators': 400}",
            1.00, 0.78],
            [ 'Decision_tree Model', r"{'criterion': 'gini', 'max_depth': 9, 'min_samples_split': 6}", 0.99, 0.59], [ '
Random forest', r"{'max_depth': 5, 'n_estimators': 400}", 1.00, 0.71],
            [ 'Calibrated Model', "--", 1.00, 0.76]])

print(t.draw())
```

| Model                  | Hyper-parameter   | Train AUC | Test AUC |
|------------------------|---|-----------|----------|
| Knn_Model              | (algorithm='kd_tree', n_neighbors=39)                         | 0.500     | 0.550    |
| logistic Regresstion   | {'C': 1, 'penalty': 'l1', 'solver': 'saga'}                   | 1         | 0.770    |
| Support Vector Machine | {'C': 1, 'kernel': 'poly'}                                    | 1         | 0.640    |
| XGBoost Classifier     | {'max_depth': 2, 'n_estimators': 400}                         | 1         | 0.780    |
| Decision_tree Model    | {'criterion': 'gini', 'max_depth': 9, 'min_samples_split': 6} | 0.990     | 0.590    |
| Random forest          | {'max_depth': 5, 'n_estimators': 400}                         | 1         | 0.710    |
| Calibrated Model       | --  | 1         | 0.760    |

We have used Calibrated Classifier and got training AUC = 1.0 and test AUC = 0.76

As per observation, XGboost is giving Good Accuracy from above applied algorithm